mlp | 96p96c4g | numa | collective vs p2p1 | rep 1 | call 6 / MAIA step 189

Overall MAIA ML-step makespan: collective 264.76 ms | p2p1 257.49 ms

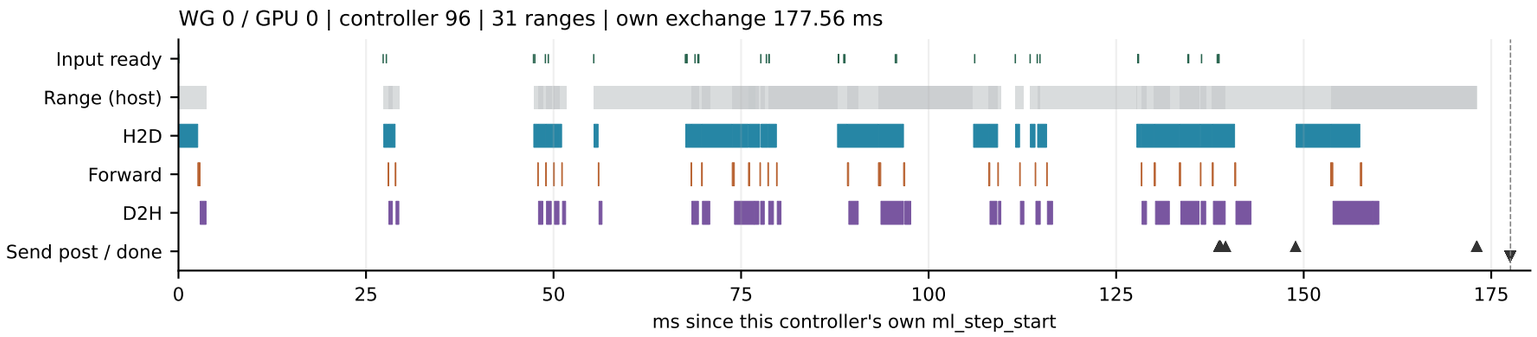

GPU work (ms)	Collective	p2p1
H2D	26.546	55.178
Forward	2.606	4.557
D2H	13.454	34.315
GPU SUM	42.606	94.049

Controllers: coll 96 / P2P 96
Coll host device elapsed: 42.762 ms
P2P GPU envelope: 159.886 ms
Different scopes; not an elapsed-time comparison.

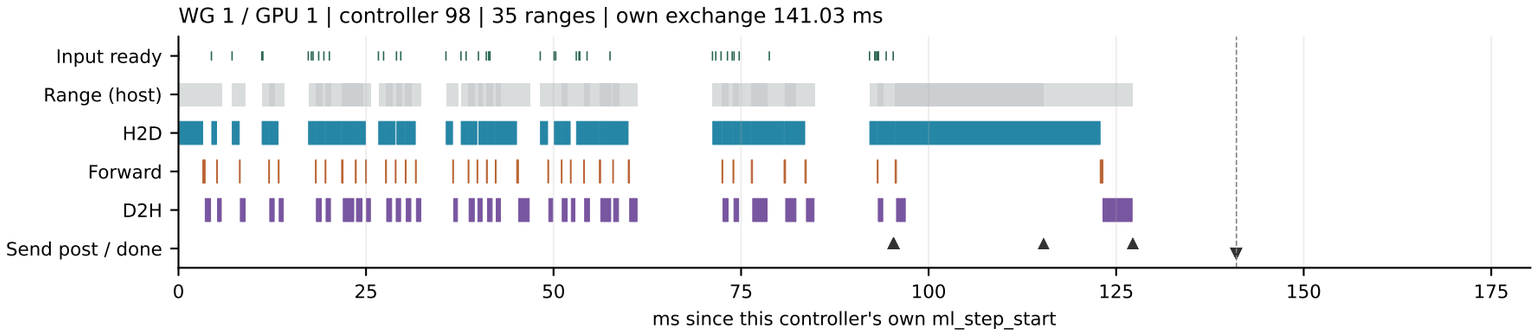

GPU work (ms)	Collective	p2p1
H2D	25.584	78.721
Forward	2.595	5.076
D2H	13.109	30.488
GPU SUM	41.288	114.285

Controllers: coll 98 / P2P 98
Coll host device elapsed: 41.468 ms
P2P GPU envelope: 126.984 ms
Different scopes; not an elapsed-time comparison.

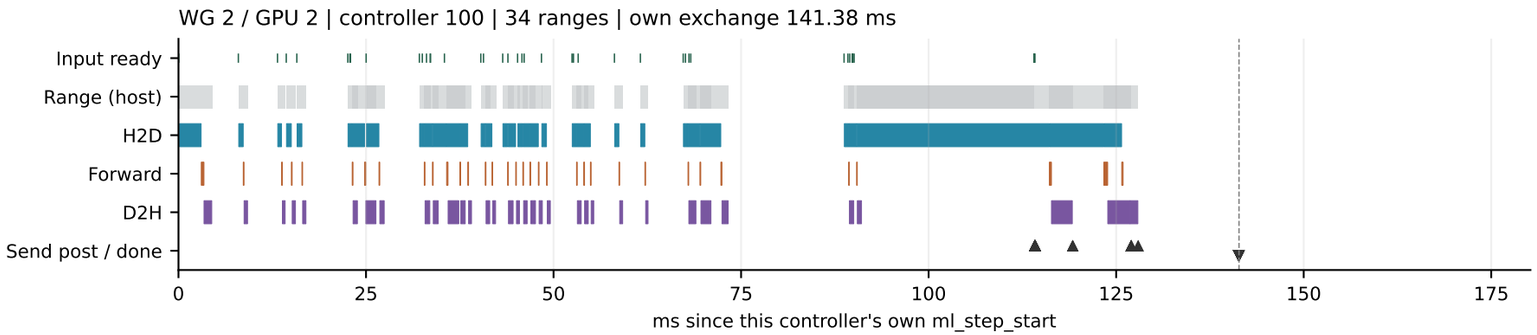

GPU work (ms)	Collective	p2p1
H2D	26.591	66.740
Forward	2.614	4.536
D2H	13.412	25.834
GPU SUM	42.618	97.109

Controllers: coll 100 / P2P 100
Coll host device elapsed: 42.806 ms
P2P GPU envelope: 127.753 ms
Different scopes; not an elapsed-time comparison.

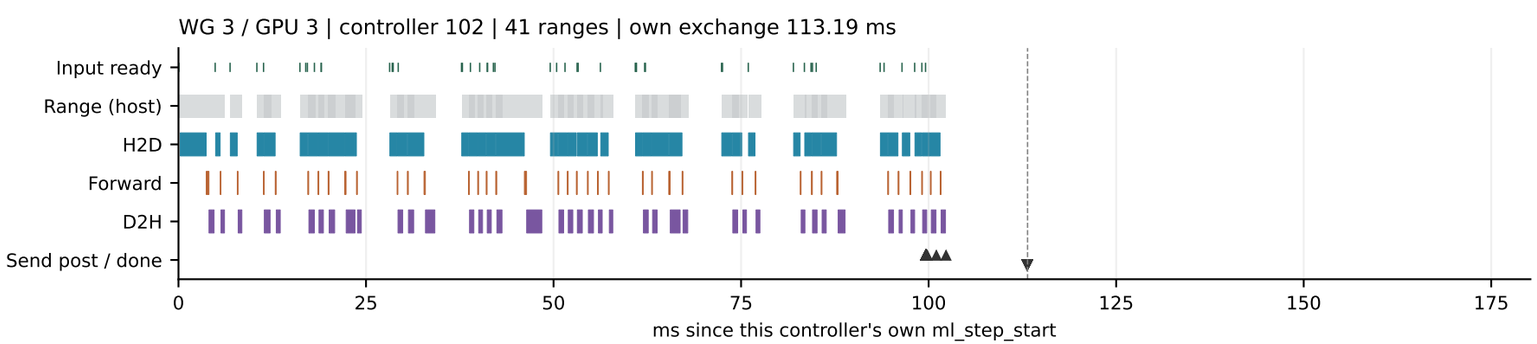

GPU work (ms)	Collective	p2p1
H2D	25.855	54.966
Forward	2.592	5.399
D2H	12.905	27.909
GPU SUM	41.352	88.274

Controllers: coll 102 / P2P 102
Coll host device elapsed: 41.534 ms
P2P GPU envelope: 102.008 ms
Different scopes; not an elapsed-time comparison.

Each controller has its own zero; rows are NOT globally synchronized. CUDA-event durations are reliable; host-anchored inter-request overlap is approximate. GPU work SUM is not elapsed time or a critical path.