transformer | 96c4g | default | collective vs p2pfull | rep 1 | call 6 / MAIA step 189

Overall MAIA ML-step makespan: collective 1719.48 ms | p2pfull 2077.90 ms

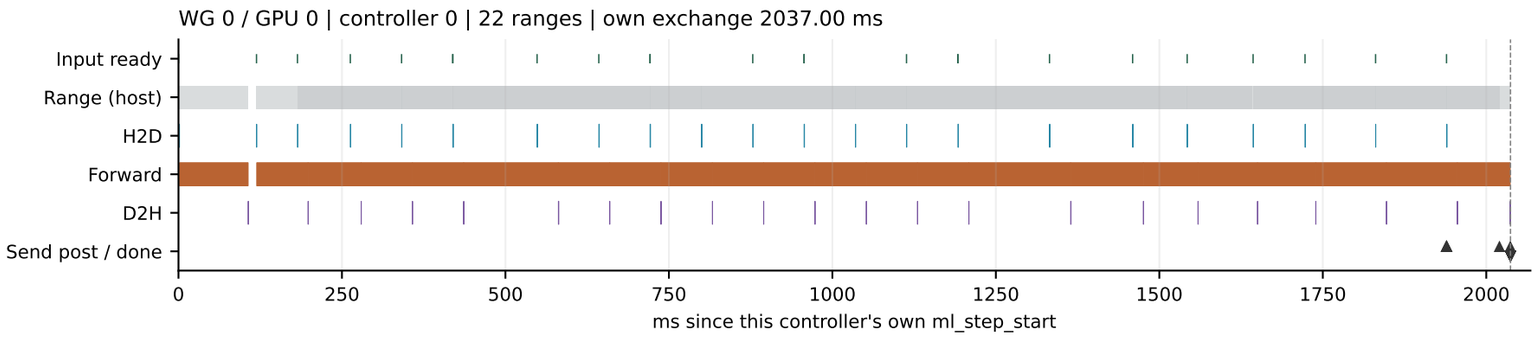

GPU work (ms)	Collective	p2pfull
H2D	5.105	12.725
Forward	1615.634	2021.840
D2H	3.652	6.017
GPU SUM	1624.391	2040.583

Controllers: coll 0 / P2P 0
Coll host device elapsed: 1624.905 ms
P2P GPU envelope: 2036.549 ms
Different scopes; not an elapsed-time comparison.

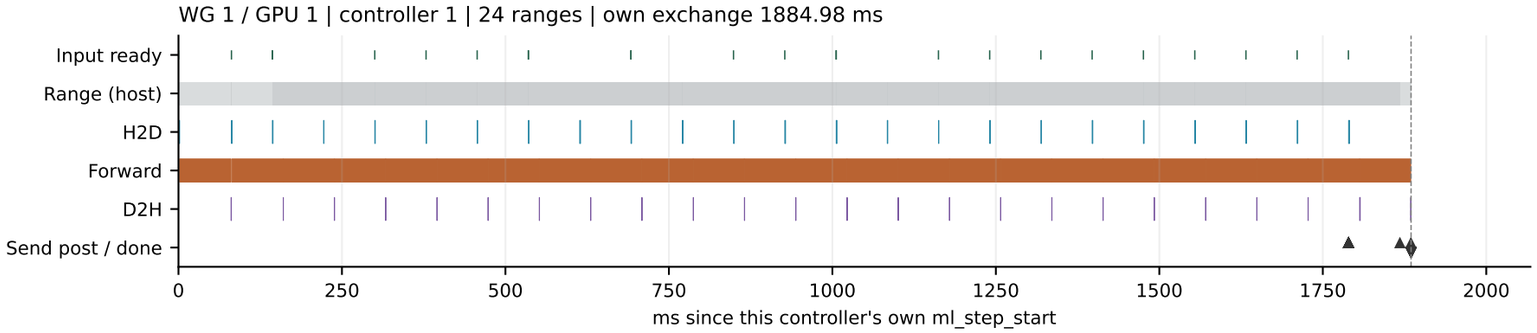

GPU work (ms)	Collective	p2pfull
H2D	9.100	22.788
Forward	1615.644	1881.205
D2H	5.282	7.871
GPU SUM	1630.026	1911.864

Controllers: coll 1 / P2P 1
Coll host device elapsed: 1630.507 ms
P2P GPU envelope: 1884.441 ms
Different scopes; not an elapsed-time comparison.

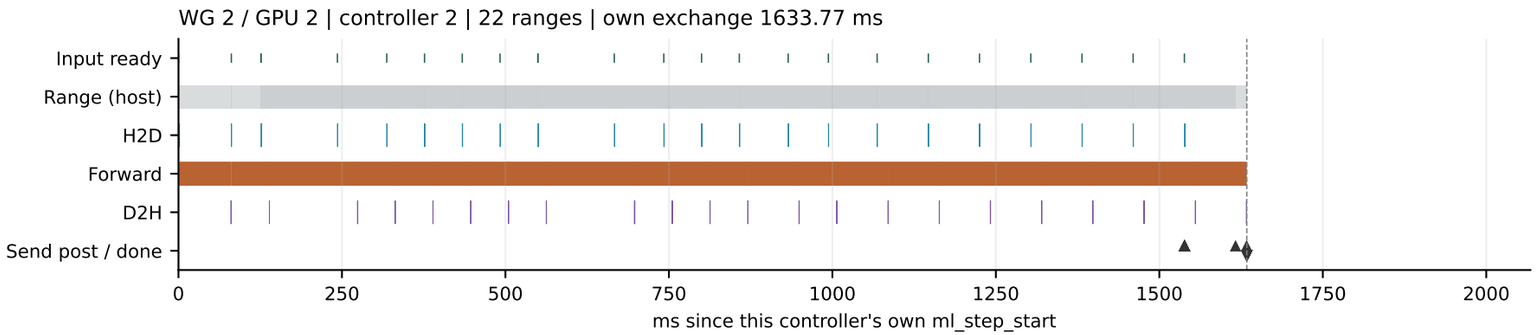

GPU work (ms)	Collective	p2pfull
H2D	4.016	12.396
Forward	1436.373	1630.317
D2H	9.097	8.495
GPU SUM	1449.485	1651.209

Controllers: coll 2 / P2P 2
Coll host device elapsed: 1450.008 ms
P2P GPU envelope: 1633.202 ms
Different scopes; not an elapsed-time comparison.

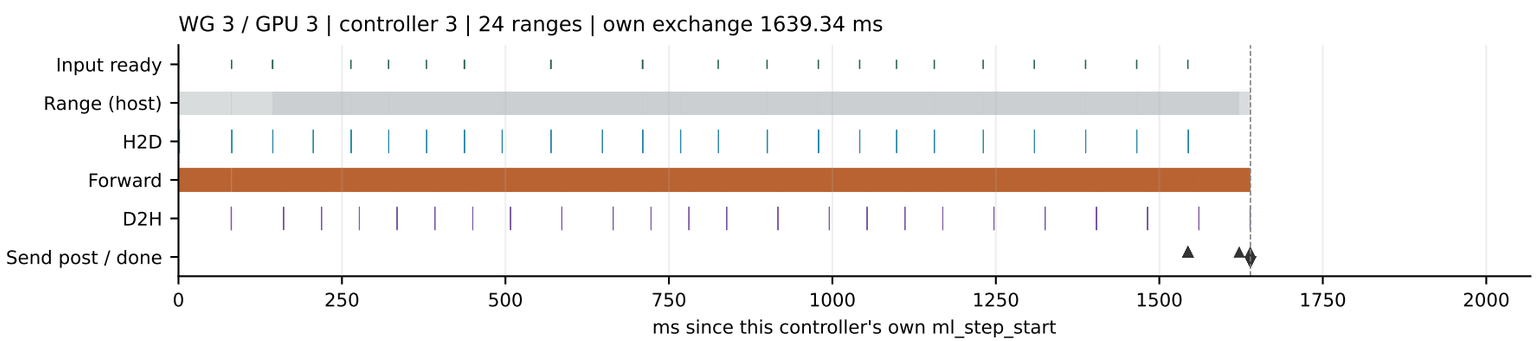

GPU work (ms)	Collective	p2pfull
H2D	4.629	12.277
Forward	1426.385	1635.469
D2H	9.514	8.325
GPU SUM	1440.527	1656.071

Controllers: coll 3 / P2P 3
Coll host device elapsed: 1441.077 ms
P2P GPU envelope: 1638.819 ms
Different scopes; not an elapsed-time comparison.

Each controller has its own zero; rows are NOT globally synchronized. CUDA-event durations are reliable; host-anchored inter-request overlap is approximate. GPU work SUM is not elapsed time or a critical path.