transformer | 96c4g | numa | collective vs p2pfull | rep 1 | call 6 / MAIA step 189

Overall MAIA ML-step makespan: collective 1693.84 ms | p2pfull 2025.05 ms

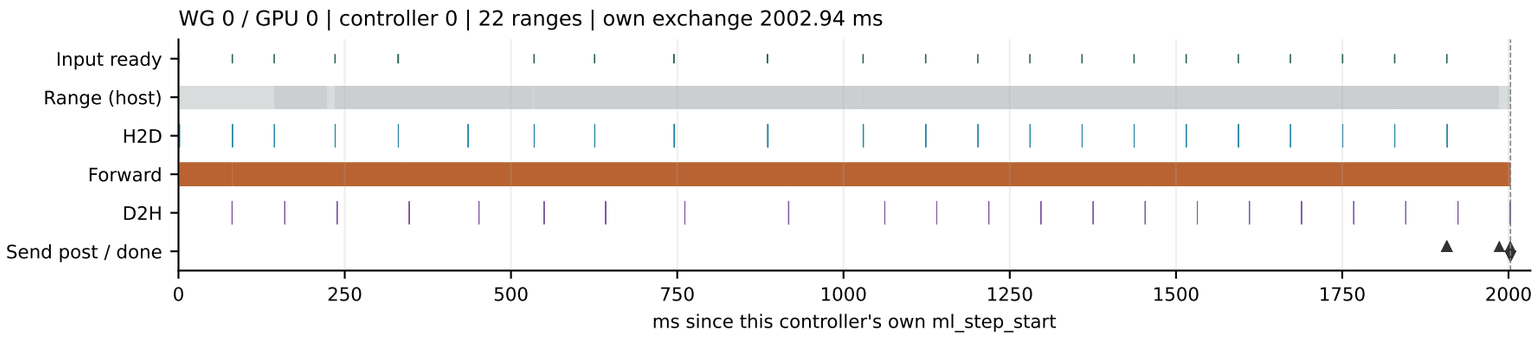

GPU work (ms)	Collective	p2pfull
H2D	4.967	11.196
Forward	1615.280	1999.466
D2H	7.014	4.592
GPU SUM	1627.261	2015.254

Controllers: coll 0 / P2P 0
Coll host device elapsed: 1627.799 ms
P2P GPU envelope: 2002.499 ms
Different scopes; not an elapsed-time comparison.

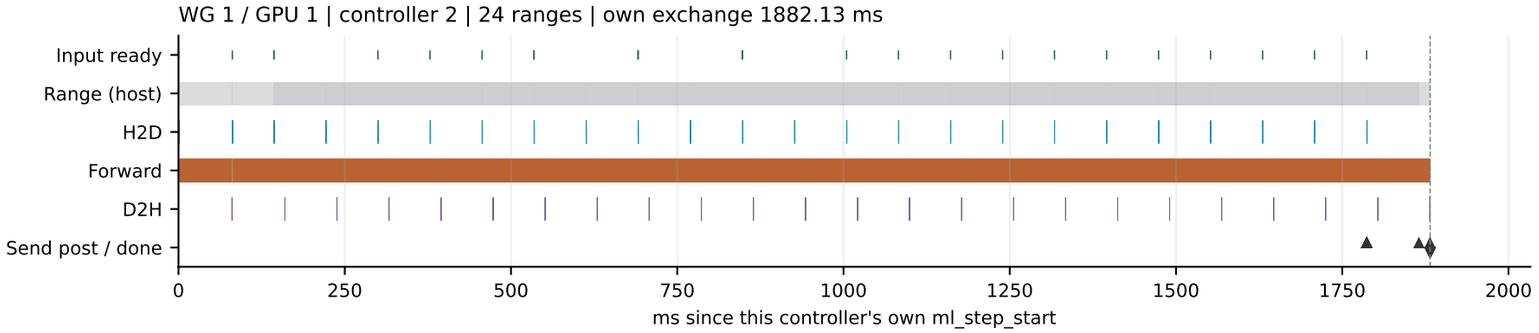

GPU work (ms)	Collective	p2pfull
H2D	4.748	12.538
Forward	1616.324	1878.752
D2H	6.568	4.670
GPU SUM	1627.640	1895.959

Controllers: coll 2 / P2P 2
Coll host device elapsed: 1628.147 ms
P2P GPU envelope: 1881.676 ms
Different scopes; not an elapsed-time comparison.

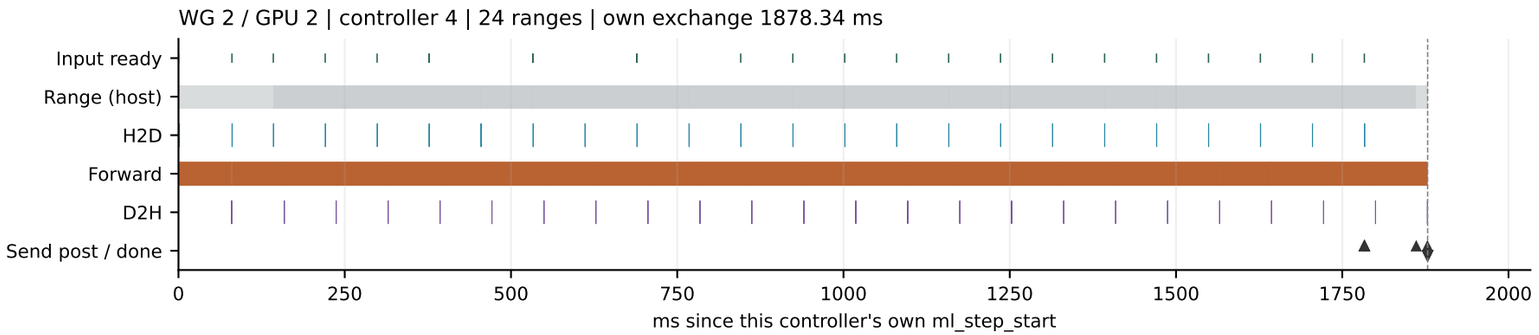

GPU work (ms)	Collective	p2pfull
H2D	4.957	11.093
Forward	1618.055	1875.168
D2H	3.915	4.681
GPU SUM	1626.928	1890.942

Controllers: coll 4 / P2P 4
Coll host device elapsed: 1627.424 ms
P2P GPU envelope: 1877.922 ms
Different scopes; not an elapsed-time comparison.

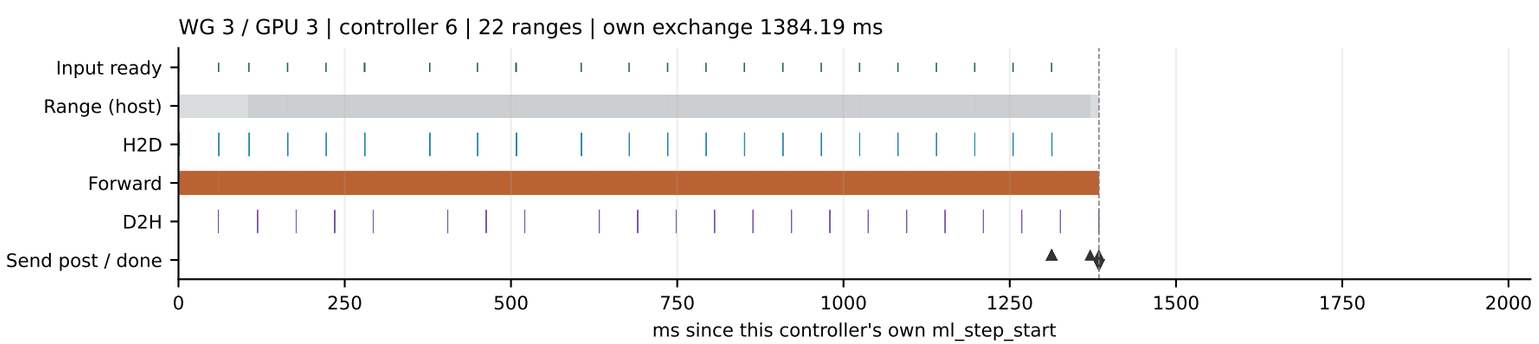

GPU work (ms)	Collective	p2pfull
H2D	5.457	10.727
Forward	1231.899	1381.380
D2H	5.395	3.723
GPU SUM	1242.751	1395.830

Controllers: coll 6 / P2P 6
Coll host device elapsed: 1243.356 ms
P2P GPU envelope: 1383.701 ms
Different scopes; not an elapsed-time comparison.

Each controller has its own zero; rows are NOT globally synchronized. CUDA-event durations are reliable; host-anchored inter-request overlap is approximate. GPU work SUM is not elapsed time or a critical path.