transformer | 96p96c4g | default | collective vs p2p1 | rep 1 | call 6 / MAIA step 189

Overall MAIA ML-step makespan: collective 1711.11 ms | p2p1 2519.17 ms

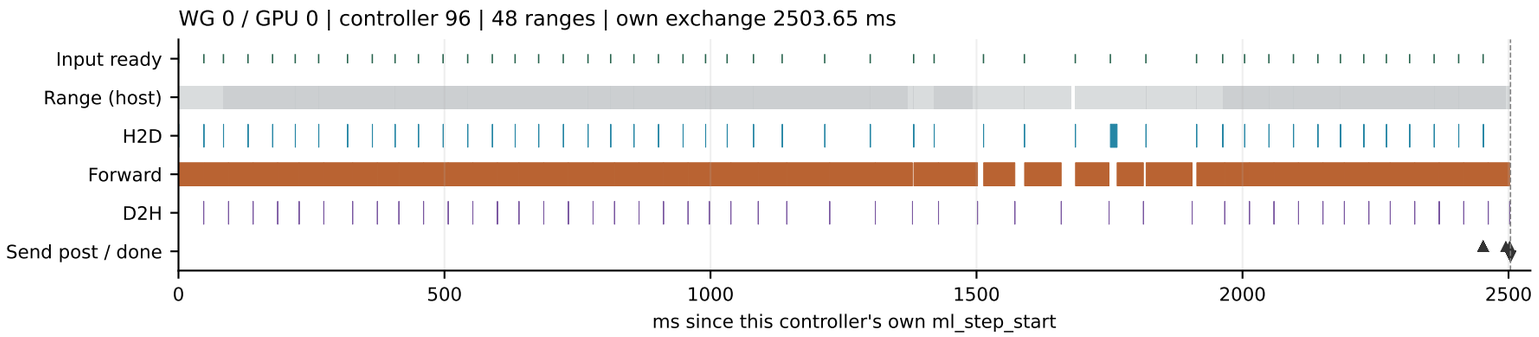

GPU work (ms)	Collective	p2p1
H2D	5.470	27.541
Forward	1644.372	2411.861
D2H	6.603	7.720
GPU SUM	1656.445	2447.122

Controllers: coll 96 / P2P 96
Coll host device elapsed: 1657.137 ms
P2P GPU envelope: 2502.041 ms
Different scopes; not an elapsed-time comparison.

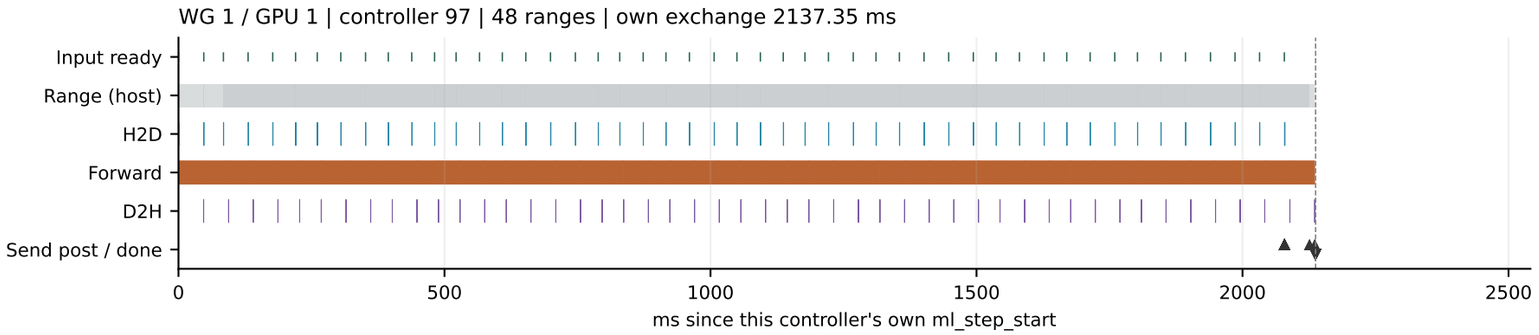

GPU work (ms)	Collective	p2p1
H2D	9.553	26.765
Forward	1617.555	2133.752
D2H	8.769	9.144
GPU SUM	1635.877	2169.660

Controllers: coll 97 / P2P 97
Coll host device elapsed: 1636.379 ms
P2P GPU envelope: 2135.608 ms
Different scopes; not an elapsed-time comparison.

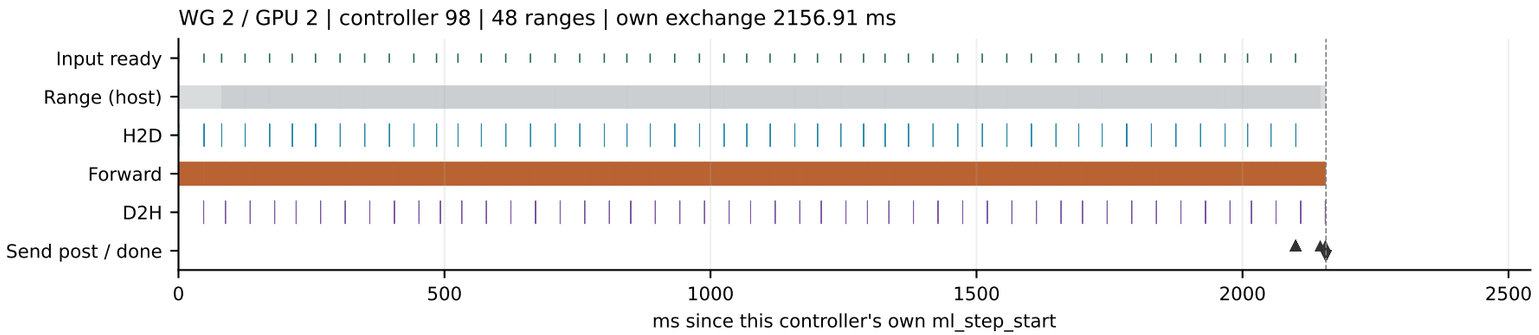

GPU work (ms)	Collective	p2p1
H2D	4.629	16.590
Forward	1652.313	2153.574
D2H	8.770	10.569
GPU SUM	1665.712	2180.733

Controllers: coll 98 / P2P 98
Coll host device elapsed: 1666.255 ms
P2P GPU envelope: 2155.564 ms
Different scopes; not an elapsed-time comparison.

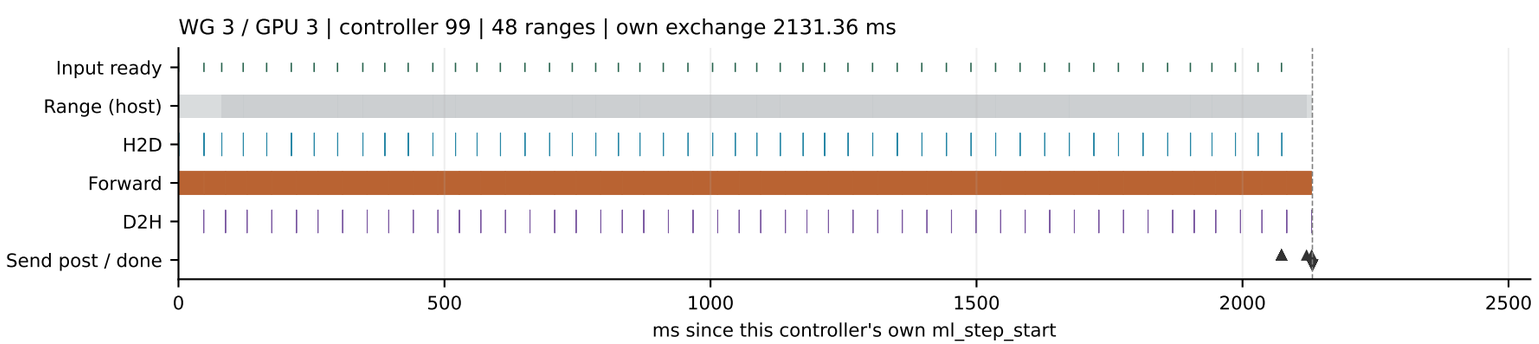

GPU work (ms)	Collective	p2p1
H2D	4.545	16.529
Forward	1615.441	2127.873
D2H	10.823	11.166
GPU SUM	1630.809	2155.568

Controllers: coll 99 / P2P 99
Coll host device elapsed: 1631.351 ms
P2P GPU envelope: 2129.673 ms
Different scopes; not an elapsed-time comparison.

Each controller has its own zero; rows are NOT globally synchronized. CUDA-event durations are reliable; host-anchored inter-request overlap is approximate. GPU work SUM is not elapsed time or a critical path.