transformer | 96p96c4g | default | collective vs p2pfull | rep 1 | call 6 / MAIA step 189

Overall MAIA ML-step makespan: collective 1711.11 ms | p2pfull 2543.07 ms

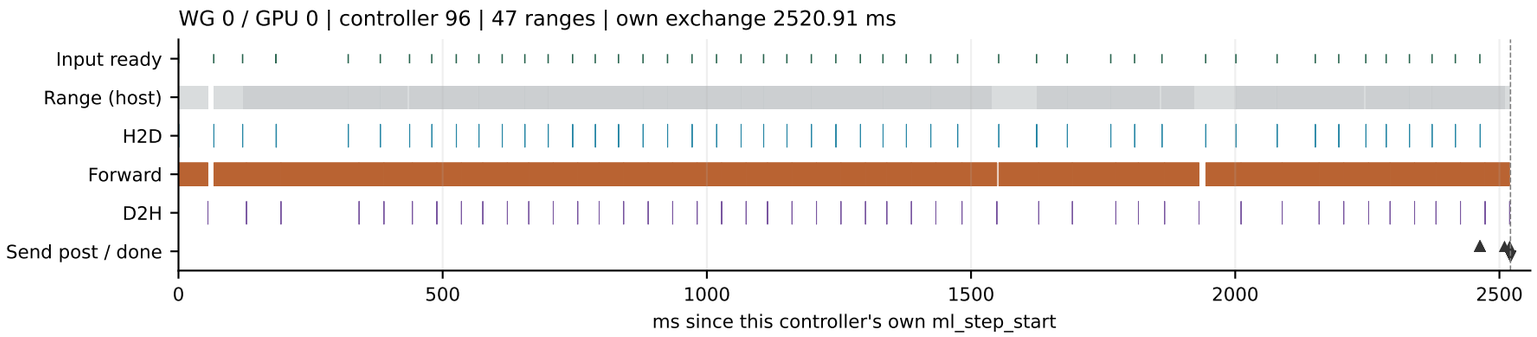

GPU work (ms)	Collective	p2pfull
H2D	5.470	16.122
Forward	1644.372	2489.819
D2H	6.603	7.770
GPU SUM	1656.445	2513.710

Controllers: coll 96 / P2P 96
Coll host device elapsed: 1657.137 ms
P2P GPU envelope: 2519.189 ms
Different scopes; not an elapsed-time comparison.

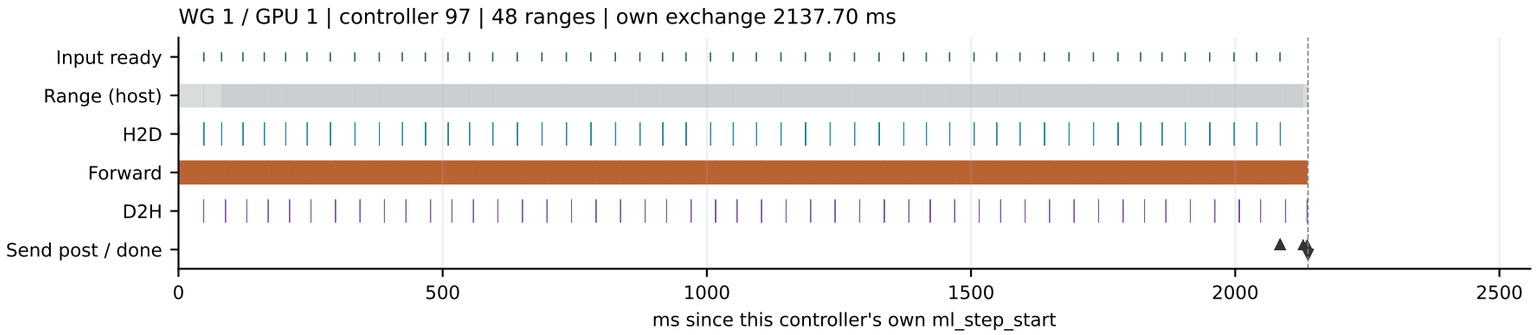

GPU work (ms)	Collective	p2pfull
H2D	9.553	26.739
Forward	1617.555	2133.988
D2H	8.769	9.026
GPU SUM	1635.877	2169.753

Controllers: coll 97 / P2P 97
Coll host device elapsed: 1636.379 ms
P2P GPU envelope: 2135.893 ms
Different scopes; not an elapsed-time comparison.

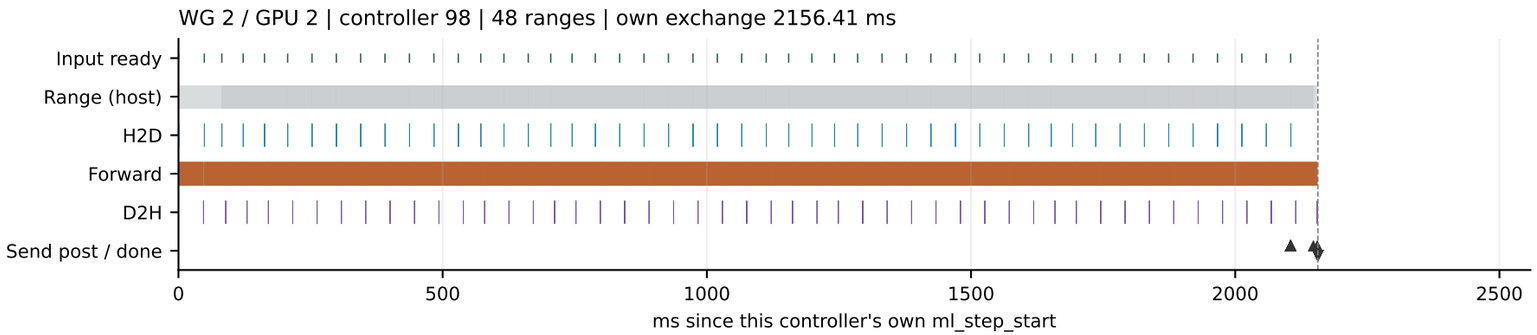

GPU work (ms)	Collective	p2pfull
H2D	4.629	14.890
Forward	1652.313	2152.164
D2H	8.770	10.759
GPU SUM	1665.712	2177.814

Controllers: coll 98 / P2P 98
Coll host device elapsed: 1666.255 ms
P2P GPU envelope: 2155.031 ms
Different scopes; not an elapsed-time comparison.

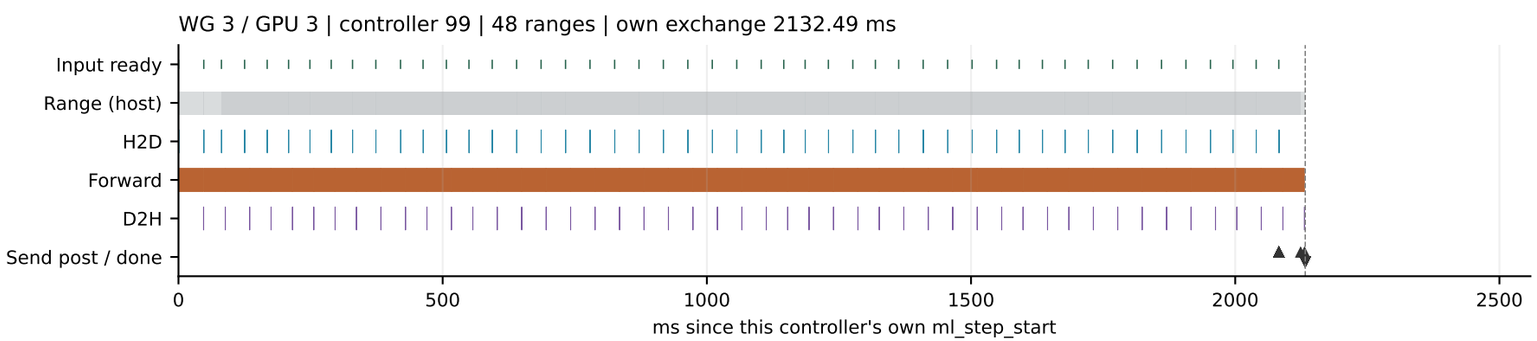

GPU work (ms)	Collective	p2pfull
H2D	4.545	14.348
Forward	1615.441	2128.916
D2H	10.823	11.225
GPU SUM	1630.809	2154.489

Controllers: coll 99 / P2P 99
Coll host device elapsed: 1631.351 ms
P2P GPU envelope: 2130.822 ms
Different scopes; not an elapsed-time comparison.

Each controller has its own zero; rows are NOT globally synchronized. CUDA-event durations are reliable; host-anchored inter-request overlap is approximate. GPU work SUM is not elapsed time or a critical path.