mlp | 24c1g | default | collective vs p2p1 | rep 1 | call 6 / MAIA step 189

Overall MAIA ML-step makespan: collective 1077.20 ms | p2p1 986.07 ms

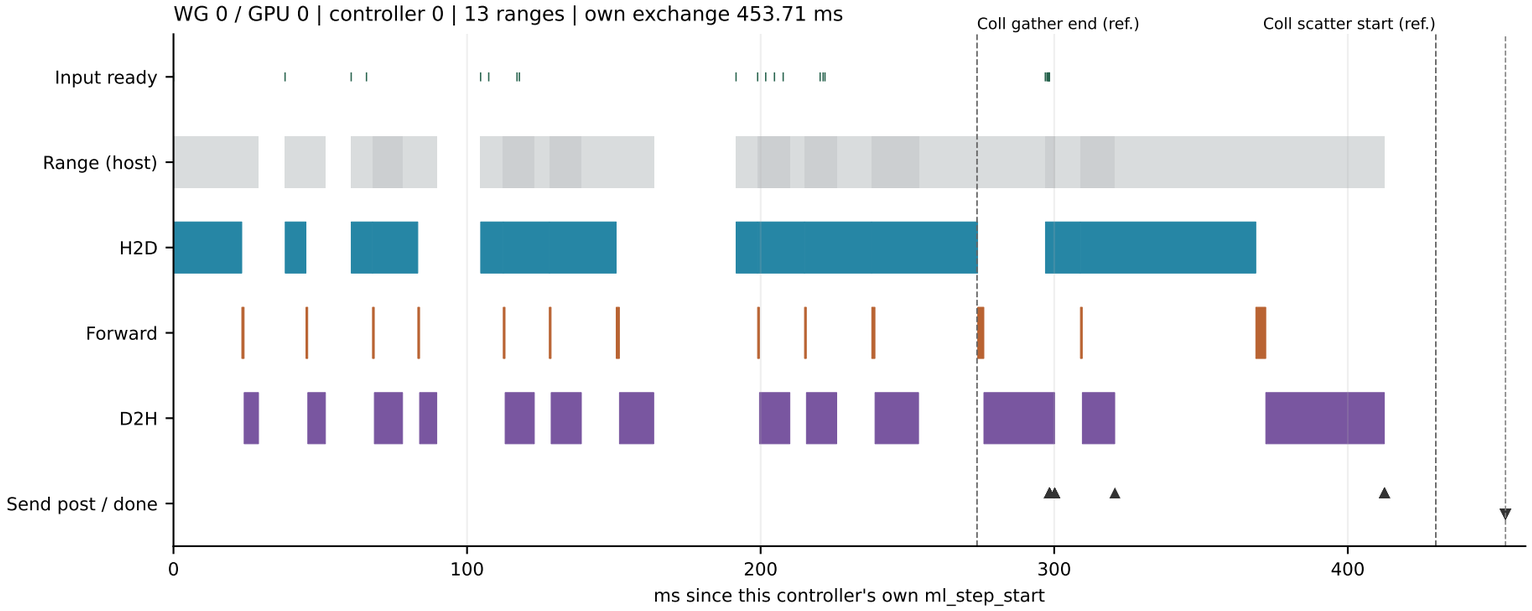

GPU work (ms)	Collective	p2p1
H2D	95.904	252.701
Forward	11.198	13.402
D2H	48.961	169.334
GPU SUM	156.063	435.437

Controllers: coll 0 / P2P 0
Coll host device elapsed: 156.273 ms
P2P GPU envelope: 412.349 ms
Different scopes; not an elapsed-time comparison.

Each controller has its own zero; rows are NOT globally synchronized. CUDA-event durations are reliable; host-anchored inter-request overlap is approximate. GPU work SUM is not elapsed time or a critical path.

Dashed collective boundaries: concatenated host durations from independent zero; gaps omitted.