# mlp | 24c1g | numa | collective vs p2p1 | rep 1 | call 6 / MAIA step 189

Overall MAIA ML-step makespan: collective 1103.16 ms | p2p1 967.22 ms

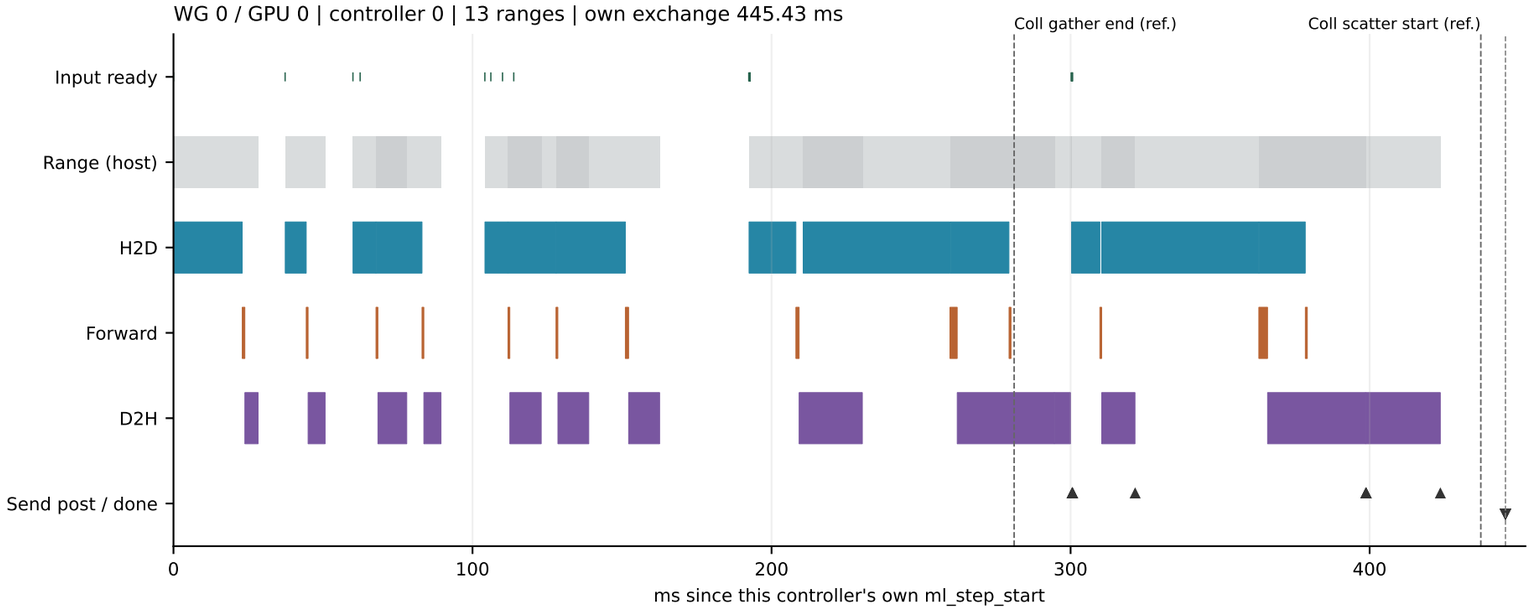


| GPU work (ms) | Collective | p2p1 |
|---|---|---|
| H2D | 95.775 | 261.552 |
| Forward | 11.205 | 13.361 |
| D2H | 48.950 | 184.502 |
| GPU SUM | 155.930 | 459.415 |

Controllers: coll 0 / P2P 0
Coll host device elapsed: 156.063 ms
P2P GPU envelope: 423.478 ms
Different scopes; not an elapsed-time comparison.

Each controller has its own zero; rows are NOT globally synchronized. CUDA-event durations are reliable; host-anchored inter-request overlap is approximate. GPU work SUM is not elapsed time or a critical path.

Dashed collective boundaries: concatenated host durations from independent zero; gaps omitted.