transformer | 24c1g | default | collective vs p2p1 | rep 1 | call 6 / MAIA step 189

Overall MAIA ML-step makespan: collective 5744.07 ms | p2p1 5890.46 ms

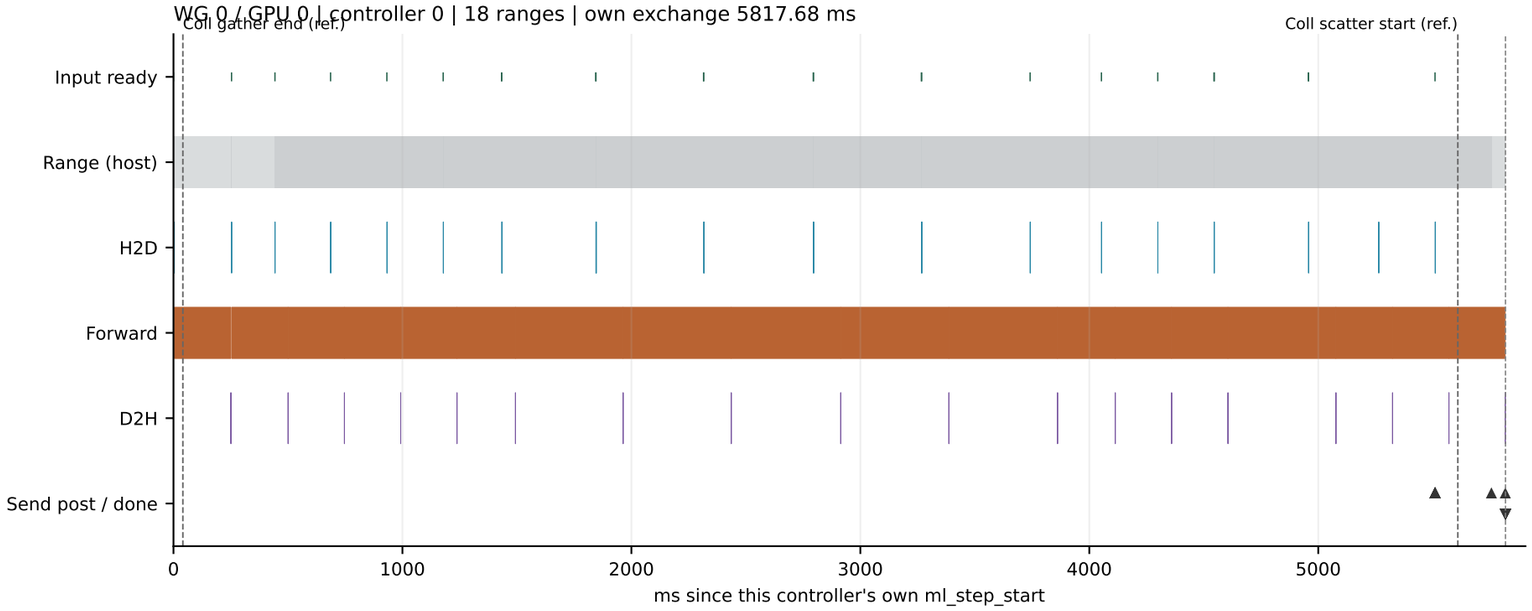

GPU work (ms)	Collective	p2p1
H2D	15.372	26.806
Forward	5555.673	5808.749
D2H	7.978	9.149
GPU SUM	5579.024	5844.703

Controllers: coll 0 / P2P 0
Coll host device elapsed: 5568.597 ms
P2P GPU envelope: 5816.814 ms
Different scopes; not an elapsed-time comparison.

Each controller has its own zero; rows are NOT globally synchronized. CUDA-event durations are reliable; host-anchored inter-request overlap is approximate. GPU work SUM is not elapsed time or a critical path.

Dashed collective boundaries: concatenated host durations from independent zero; gaps omitted.