mlp | 24c1g | numa | collective vs p2pfull | rep 1 | call 6 / MAIA step 189

Overall MAIA ML-step makespan: collective 1103.16 ms | p2pfull 1010.40 ms

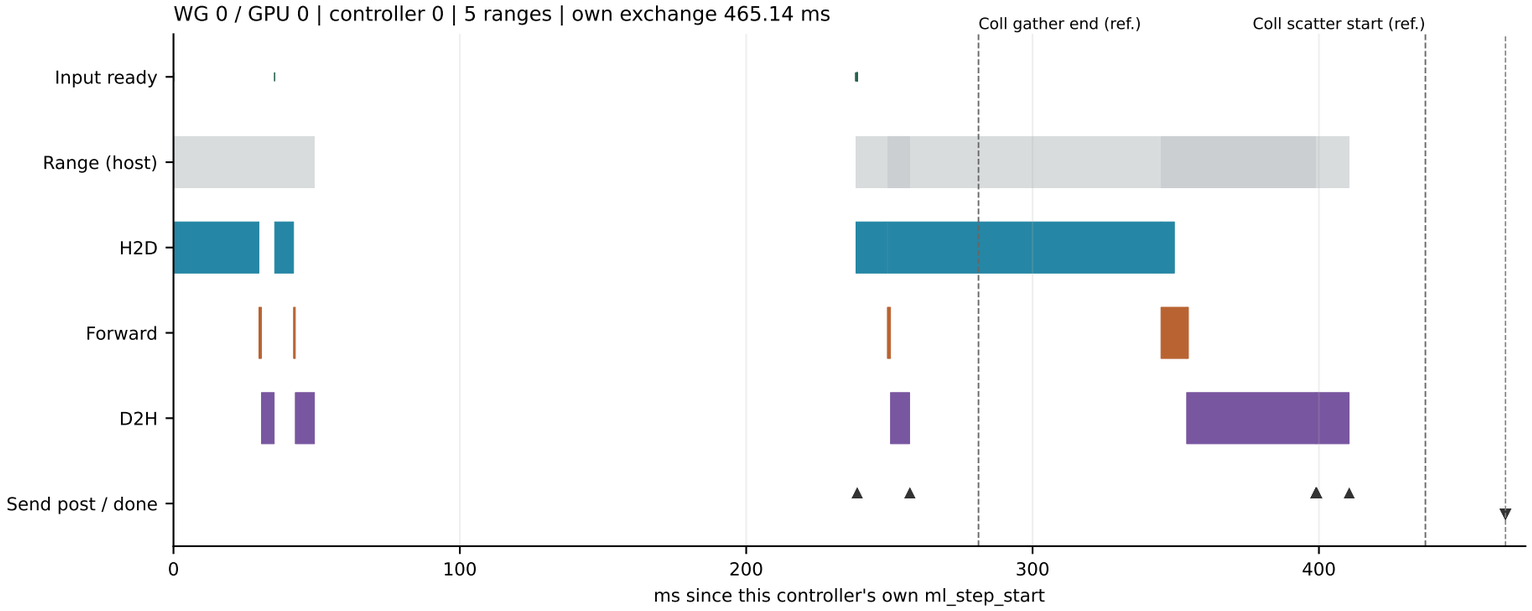

GPU work (ms)	Collective	p2pfull
H2D	95.775	147.056
Forward	11.205	12.141
D2H	48.950	74.526
GPU SUM	155.930	233.723

Controllers: coll 0 / P2P 0
Coll host device elapsed: 156.063 ms
P2P GPU envelope: 410.305 ms
Different scopes; not an elapsed-time comparison.

Each controller has its own zero; rows are NOT globally synchronized. CUDA-event durations are reliable; host-anchored inter-request overlap is approximate. GPU work SUM is not elapsed time or a critical path.

Dashed collective boundaries: concatenated host durations from independent zero; gaps omitted.