transformer | 24c1g | default | collective vs p2pfull | rep 1 | call 6 / MAIA step 189

Overall MAIA ML-step makespan: collective 5744.07 ms | p2pfull 5883.09 ms

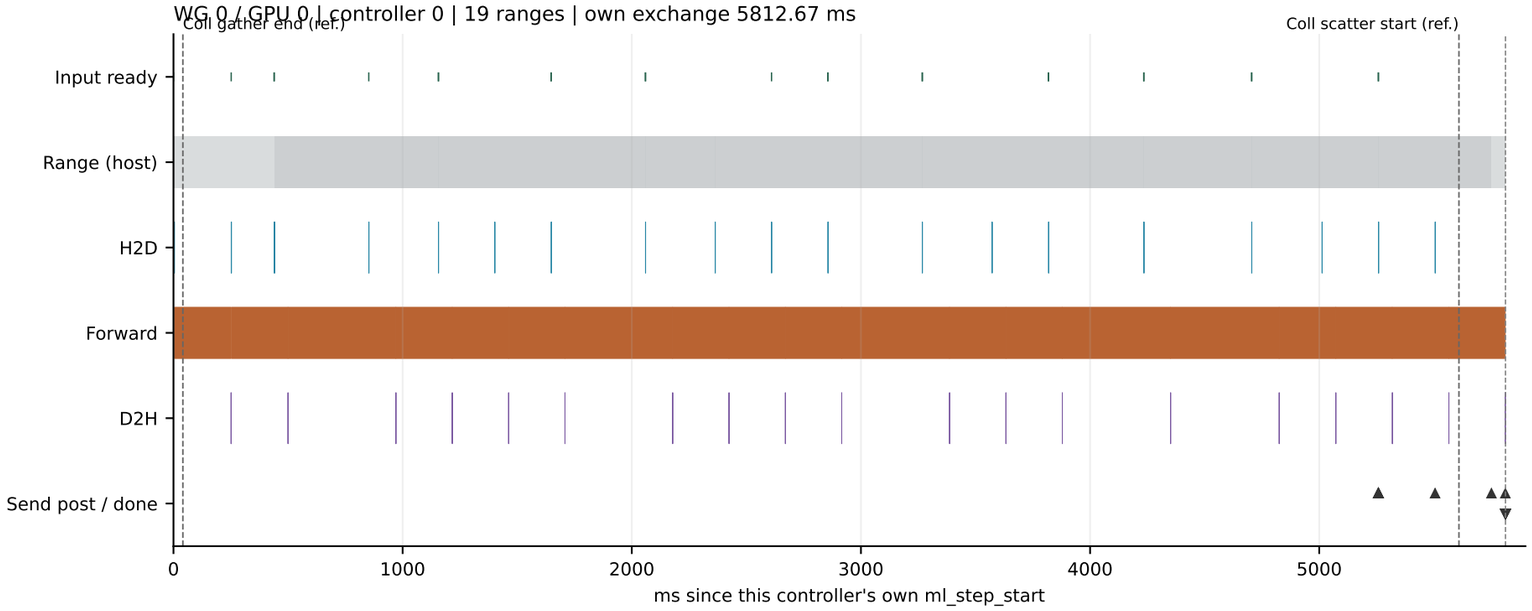

GPU work (ms)	Collective	p2pfull
H2D	15.372	23.223
Forward	5555.673	5804.848
D2H	7.978	9.470
GPU SUM	5579.024	5837.541

Controllers: coll 0 / P2P 0
Coll host device elapsed: 5568.597 ms
P2P GPU envelope: 5811.880 ms
Different scopes; not an elapsed-time comparison.

Each controller has its own zero; rows are NOT globally synchronized. CUDA-event durations are reliable; host-anchored inter-request overlap is approximate. GPU work SUM is not elapsed time or a critical path.

Dashed collective boundaries: concatenated host durations from independent zero; gaps omitted.