transformer | 24c1g | numa | collective vs p2p1 | rep 1 | call 6 / MAIA step 189

Overall MAIA ML-step makespan: collective 5737.07 ms | p2p1 5881.42 ms

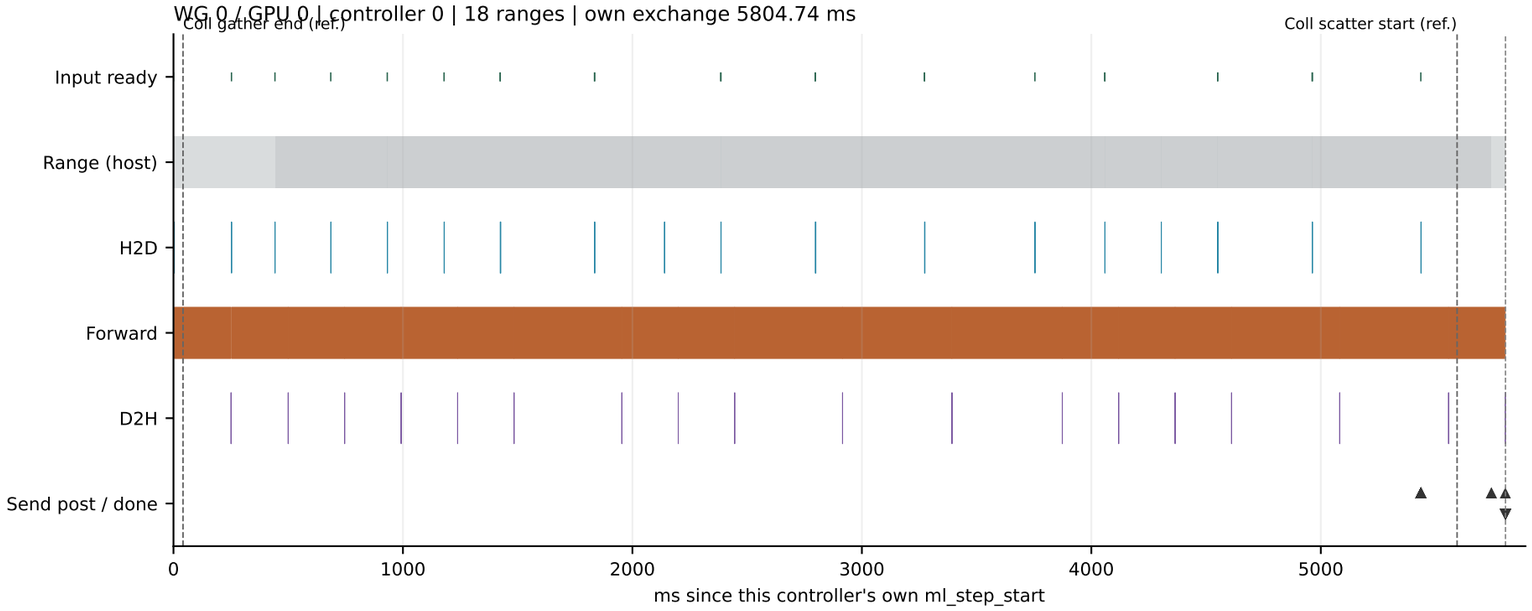

GPU work (ms)	Collective	p2p1
H2D	15.389	25.753
Forward	5539.443	5796.824
D2H	8.002	8.589
GPU SUM	5562.834	5831.166

Controllers: coll 0 / P2P 0
Coll host device elapsed: 5552.329 ms
P2P GPU envelope: 5803.999 ms
Different scopes; not an elapsed-time comparison.

Each controller has its own zero; rows are NOT globally synchronized. CUDA-event durations are reliable; host-anchored inter-request overlap is approximate. GPU work SUM is not elapsed time or a critical path.

Dashed collective boundaries: concatenated host durations from independent zero; gaps omitted.