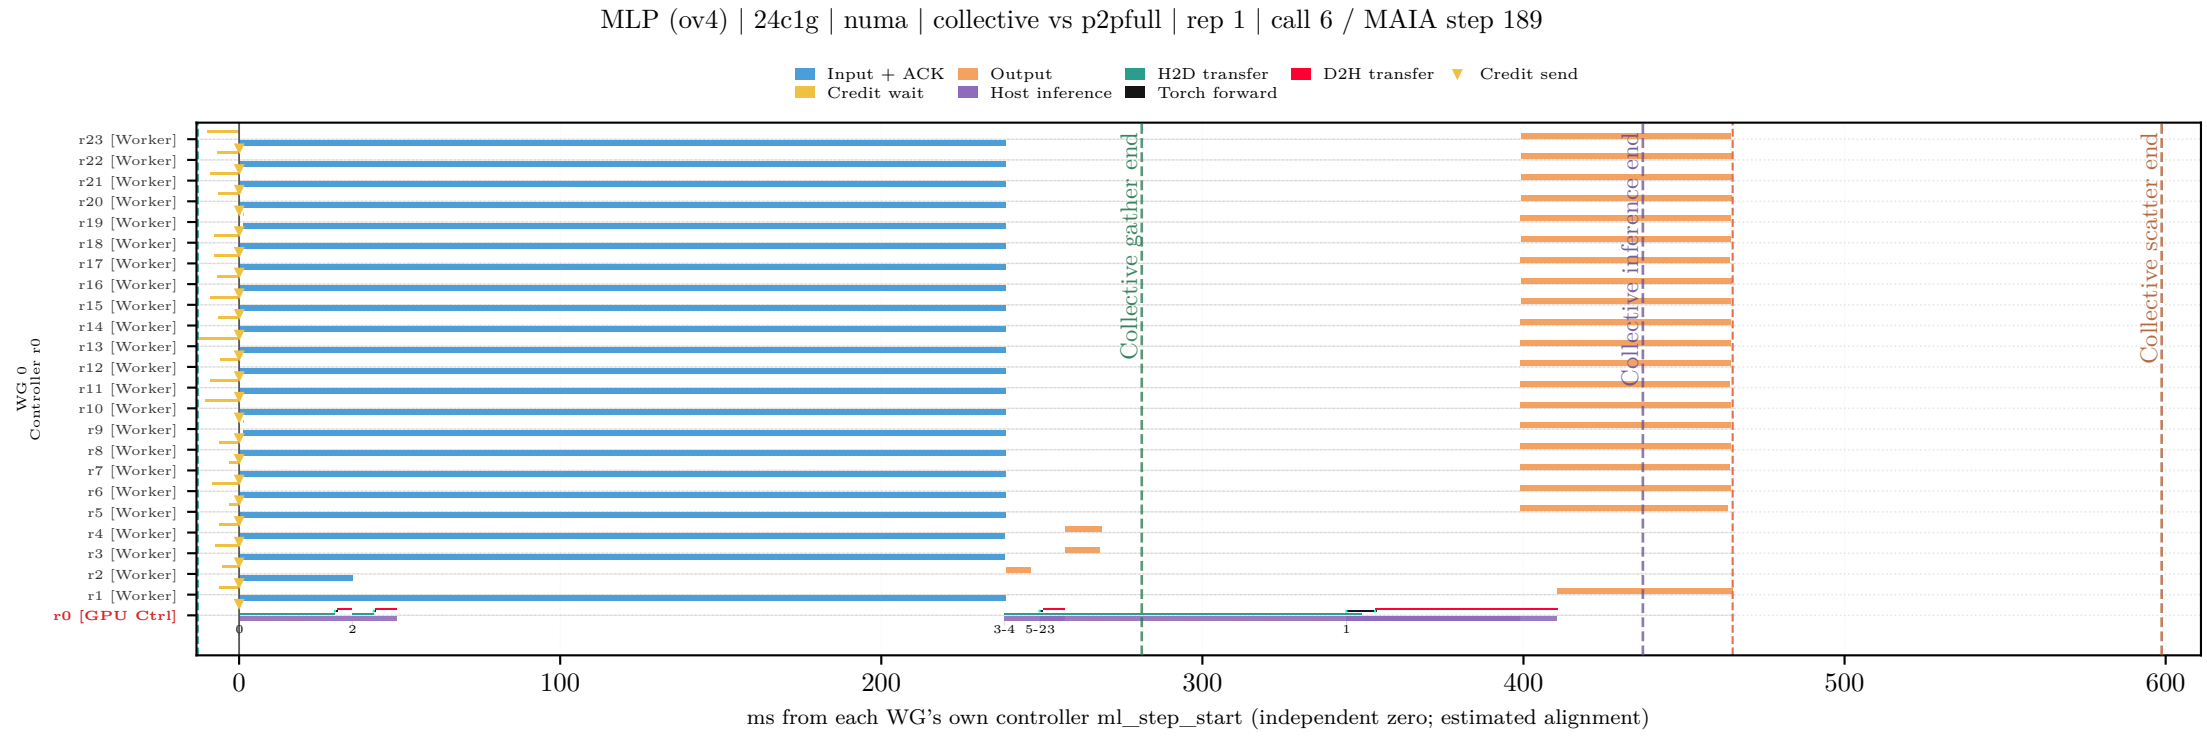

MLP (ov4) | 24c1g | numa | collective vs p2pfull | rep 1 | call 6 / MAIA step 189
Input + ACK
Credit wait
Output
Host inference
H2D transfer
Torch forward
D2H transfer
Credit send
WG 0
Controller r0
r23 [Worker]
r22 [Worker]
r21 [Worker]
r20 [Worker]
r19 [Worker]
r18 [Worker]
r17 [Worker]
r16 [Worker]
r15 [Worker]
r14 [Worker]
r13 [Worker]
r12 [Worker]
r11 [Worker]
r10 [Worker]
r9 [Worker]
r8 [Worker]
r7 [Worker]
r6 [Worker]
r5 [Worker]
r4 [Worker]
r3 [Worker]
r2 [Worker]
r1 [Worker]
r0 [GPU Ctrl]
Collective gather end
Collective inference end
Collective scatter end
0
2
3-4
5-23
1
0
100
200
300
400
500
600
ms from each WG's own controller ml_step_start (independent zero; estimated alignment)